

base LLM: Gemma-4-12B, Method: PreSTO-EntropyCut, inference engine: vLLM

 ϵ -certain reject  ϵ -uncertain  ϵ -certain accept

