

PreSTO-EntropyCut (accept first; best $N_{\text{pf}} = 20$) vs. EntropyCut

dataset: GPQA
base LLM: Qwen3.5-4B
backend: vLLM

