


(a) PreSTO-EntropyCut  
(accept first; best  $N_{\text{pf}} = 20$ )



dataset: GPQA  
base LLM: Qwen3.5-4B  
backend: vLLM

— individual trace  
— mean  
±1 sample SD

cumulative number of model calls