

dataset: GPQA, base LLM: NVIDIA-Nemotron-3-Nano-4B-BF16, inference engine: vLLM

EntropyCut — trace — mean  $\pm 1$  SD

PreSTO-EntropyCut — trace — mean  $\pm 1$  SD

