


PreSTO-EntropyCut (accept first; best $N_{\text{pf}} = 20$)
vs. EntropyCut



dataset: GPQA

base LLM: NVIDIA-Nemotron-3-Nano-4B-BF16

backend: vLLM

— PreSTO-EntropyCut trace

— PreSTO-EntropyCut mean

±1 sample SD

--- EntropyCut

cumulative number of model calls