


PreSTO-EntropyCut (accept first; best $N_{\text{pf}} = 10$)
vs. EntropyCut



dataset: GPQA

base LLM: NVIDIA-Nemotron-3-Nano-4B-BF16

backend: vLLM

- PreSTO-EntropyCut trace
- PreSTO-EntropyCut mean
- PreSTO-EntropyCut ± 1 SD
- EntropyCut trace
- EntropyCut mean
- EntropyCut ± 1 SD

cumulative model-call time (sec)