


(a) PreSTO-EntropyCut  
(accept first; best  $N_{\text{pf}} = 20$ )



dataset: GPQA

base LLM: NVIDIA-Nemotron-3-Nano-4B-BF16

backend: vLLM

— individual trace

— mean

±1 sample SD

cumulative number of model calls