


(a) PreSTO-EntropyCut
(accept first; best $N_{\text{pf}} = 20$)



dataset: GPQA
base LLM: Qwen3-4B
backend: vLLM

— individual trace
— mean
±1 sample SD

cumulative number of model calls