


PreSTO-EntropyCut (accept first; best $N_{\text{pf}} = 20$)
vs. EntropyCut



dataset: HumanEval
base LLM: NVIDIA-Nemotron-3-Nano-4B-BF16
backend: vLLM

- PreSTO-EntropyCut trace
- PreSTO-EntropyCut mean
- ±1 sample SD
- EntropyCut