

PreSTO-EntropyCut (accept first; best $N_{\text{pf}} = 10$)
vs. EntropyCut

dataset: HumanEval
base LLM: NVIDIA-Nemotron-3-Nano-4B-BF16
backend: vLLM

