


# PreSTO-EntropyCut (accept first; best $N_{\text{pf}} = 20$ ) vs. EntropyCut



dataset: HumanEval  
base LLM: Qwen3.5-4B  
backend: vLLM

- PreSTO-EntropyCut trace
- PreSTO-EntropyCut mean
- ±1 sample SD
- EntropyCut

cumulative number of model calls