


PreSTO-EntropyCut (accept first; best $N_{\text{pf}} = 20$) vs. EntropyCut



dataset: HumanEval
base LLM: Qwen3.5-4B
backend: vLLM

- PreSTO-EntropyCut trace
- PreSTO-EntropyCut mean
- PreSTO-EntropyCut ± 1 SD
- EntropyCut trace
- EntropyCut mean
- EntropyCut ± 1 SD