



dataset: LCB V6, base LLM: qwen3.5-9b




peak GPU memory delta (MB)

(a)



(b)



prefix-cache miss tokens per answer (10^3)

(c)

