


# PreSTO-EntropyCut (accept first; best $N_{\text{pf}} = 12$ ) vs. EntropyCut

dataset: LiveCodeBench-release\_v6  
base LLM: Qwen3.5-9B  
backend: vLLM



- PreSTO-EntropyCut trace
- PreSTO-EntropyCut mean
- $\pm 1$  sample SD
- EntropyCut

cumulative number of model calls