



dataset: LCB V6, base LLM: qwen3.5-9b

- PreSTO-EntropyCut (accept first)
- PreSTO-EntropyCut (reject first)
- PreSTO-EntropyCut (longest first)
- PreSTO-EntropyCut (smallest cut diff)
- PreSTO-EntropyCut (bfs accept first)
- PreSTO-EntropyCut (bfs reject first)
- EntropyCut

peak GPU memory delta (MB)



prefix-cache token hit rate (%)



prefix-cache miss tokens per answer (10^3)

