


PreSTO-EntropyCut (accept first; best $N_{\text{pf}} = 20$)
vs. EntropyCut

dataset: LiveCodeBench-release_v6
base LLM: Qwen3-8B
backend: vLLM



- PreSTO-EntropyCut trace
- PreSTO-EntropyCut mean
- ±1 sample SD
- EntropyCut

cumulative number of model calls