


(a) PreSTO-EntropyCut
(accept first; best $N_{\text{pf}} = 20$)



dataset: LiveCodeBench-release_v6
base LLM: Qwen3.5-4B
backend: vLLM

— individual trace
— mean
±1 sample SD

cumulative number of model calls