

dataset: GPQA, base LLM: Nemotron-4B, inference engine: vLLM, BFS accept-first, $B = 20$

PowerMH — trace — mean ± 1 SD

PreSTO-PowerMH — trace — mean ± 1 SD

