


SubTreeMH (accept first; best $N_{\text{pf}} = 14$) vs. PowerMH



dataset: MATH500
base LLM: Qwen2.5-7B
backend: vLLM

- SubTreeMH trace
- SubTreeMH mean
- ± 1 sample SD
- PowerMH

cumulative number of model calls