

dataset: MMLU, base LLM: Qwen3-8B, inference engine: vLLM, BFS accept-first, $B = 20$

PowerMH — trace — mean ± 1 SD

STeP-PowerMH — trace — mean ± 1 SD

